

Codex, Claude Code, and ChatGPT Pro in a Formalized Graphon Classification: An Observational Case Study

Author names to be confirmed*

Draft of September 29, 2026

Abstract

We reconstruct a six-week AI-assisted research campaign that classified a chromatic-polynomial graphon inequality for all 117 relevant graphs on at most six vertices. The resulting companion mathematics paper reports 94 positive and 23 negative rows, with all rows represented by Lean theorems. The present article asks how ChatGPT Pro, Codex, and Claude Code moved from proof search to checkable artifacts. Its evidence consists of surviving prompts and responses, session indexes, git history, mathematical notes, exact certificates, Lean sources, failure ledgers, and retained sequential build measurements. Five contrasts are supported. First, Codex did not merely translate the Atlas 43 proof: it found a shared-edge semantic error, repaired the probabilistic interpretation, reported the repair, and withheld verified status pending the remaining exact checks. Second, memory and time constraints changed proof objects: large certificates were decomposed or compressed, and Claude Code replaced an infeasible stronger route for Atlas 127 by a smaller row-level certificate. Third, formalization often narrowed a general informal theorem to the exact proposition needed for one classification row. Fourth, Codex used several proof families in a recurring order before turning to general rooted sum-of-squares certificates. Fifth, the successful long Codex run occurred in a particular operating regime—a continuous, tightly specified task with reusable exact-certificate tooling—whereas the available high, xhigh, and ultra sessions do not isolate a causal effect of reasoning effort. Because the campaign was not controlled and some source material is missing, we conclude with a prespecified clean-rerun protocol rather than a model-ranking claim.

1 Introduction

Language-model theorem proving is often evaluated as a collection of short, independent theorem attempts. Systems such as GPT-f, LeanDojo, and Baldur measure formal proof generation or repair under controlled benchmark conditions [3, 6, 2]. Other systems combine learned proposals with symbolic evaluation to search for new mathematical constructions [4, 5]. The campaign studied here has a different shape. It lasted from 13 August to 23 September 2026, used three commercial interfaces over many sessions, and ended in a repository containing mathematical notes, exact rational certificates, Lean sources, and verification scripts for a finite research classification.

The mathematical question is the following. For a finite graph H and a graphon W of edge density p , define

$$\Phi_H(p) = (1 - p)^{v(H)} \chi_H\left(\frac{1}{1 - p}\right). \quad (1)$$

*The author list and affiliations will be inserted after confirmation.

For each non-bipartite graph H on at most six vertices, excluding isolated vertices and isolated-edge components, the project asked whether

$$t(H, W) \geq \Phi_H(p) \quad \text{whenever} \quad p \geq 1 - \frac{1}{\chi(H) - 1}. \quad (2)$$

There are 117 isomorphism classes in scope. The companion mathematics paper proves that 94 satisfy (2) and 23 do not [1]. This article does not repeat those proofs. It studies the route by which the proof artifacts were created and checked.

The headline claim is not that verification is useful. Rather, the surviving materials expose several non-equivalent transitions that the word “verification” can hide: testing the semantics of an informal identity, checking exact coefficient equalities, translating a certificate into Lean, auditing theorem status, and compiling a large source tree within a stated resource envelope. Different transitions caught different problems. They also fed back into the mathematics by changing which theorem or certificate was used.

This is an observational case study. ChatGPT Pro, Codex, and Claude Code were not assigned matched tasks under fixed conditions. Model versions, reasoning effort, prompt form, repository state, session continuity, and elapsed time changed together. The human author also chose the allocation of work and the acceptable proof standard. Consequently, the paper supports claims about what happened within this campaign, but not causal comparisons among vendors, interfaces, or effort settings.

The contributions are:

1. an evidence-linked reconstruction of a long mathematical campaign, with exact quotations separated into a chronological supplement;
2. five findings based on before–after or cross-stage contrasts rather than on the mere existence of a workflow component;
3. an analysis of proof-object size as part of mathematical proof search, illustrated by Atlas 43 and Atlas 127;
4. a separation between theorem discovery, row closure, exact checking, Lean formalization, and fresh compilation; and
5. a protocol for a future clean rerun that could test hypotheses which the historical campaign alone cannot settle.

2 Case and participants

2.1 Mathematical endpoint

The campaign ended with a complete 94/23 classification. Positive rows were covered by analytic families, closure arguments, row-specific inequalities, or exact rooted sum-of-squares certificates. Every negative row received an exact rational step-graphon witness. The repository contains 117 Lean example modules whose final status fields are `.verified`; a source scan finds no `sorry` in those modules. The author additionally reports checking a successful full build on another device. The directory for that build and its exact start and end times are not present in this checkout, so this article distinguishes the author statement from the retained row-level build measurements.

The companion mathematics paper is the authority for the theorem, proof families, counterexamples, and formal statements. Here the 117 rows are units whose status changes

Table 1: AI systems in the campaign. Model identifiers are stated only where the surviving session source supplies them.

interface	observed use	surviving evidence
ChatGPT Pro	conceptual theorem search, informal derivations, and proposed proof directions	nine campaign and seven satellite conversations in the export; files and some reasoning context are absent
Codex	repository-wide classification, computational searches, proof notes, certificate infrastructure, Lean code, and status maintenance	two device indexes, including 107 workstation turns and 40 second-device primary turns; important sessions identify <code>gpt-5.6-sol</code> and their effort settings
Claude Code	later audit, certificate compaction, source compilation, and completion of the final rows	retained Claude Code sessions, including <code>claude-fable-5-1</code> and <code>claude-opus-5</code> , plus generated source and build summaries

over time: *claimed resolved* when a positive proof or negative witness was asserted, and *Lean-marked verified* when the repository status said that a Lean theorem had been completed. The second label follows historical source state; it does not itself certify that a fresh full build had already occurred.

2.2 Systems and assigned roles

Table 1 summarizes how the three interfaces appear in the surviving materials. These roles were largely assigned by the human author. Their division is therefore provenance, not a finding that the systems spontaneously specialized.

The author states that neither human collaborator wrote proof text, certificate files, scripts, or Lean code in this campaign. They supplied prompts and transferred outputs among interfaces and machines. Surviving patch-level sessions support that statement where such sessions remain, but deleted Claude Code conversations and pasted outputs prevent an archive-only proof of a universal negative. We therefore present it as an author declaration corroborated by the available traces.

3 Materials and reconstruction method

3.1 Corpus

The reconstruction uses seven kinds of material:

1. machine-readable Codex and Claude Code session indexes from two devices, with turn times, prompts, final responses, token counters, and file references;
2. a ChatGPT export extract and a manually reconstructed prompt appendix;
3. git commits, trailers, and source timestamps;
4. the live 117-row catalogue and the numbered failed-route ledger;
5. mathematical notes and note-to-Lean correspondence sections;
6. exact certificate and counterexample checkers; and
7. retained per-row and Atlas 43 sequential-build summaries.




Figure 1: Catalogue progress against an incomplete, non-additive effort proxy. The blue curve follows claimed mathematical resolution; the orange curve follows historical Lean status. Neither curve should be interpreted as a controlled learning curve.

We classify a statement as *logged* when it occurs in a surviving prompt or response, *source-level* when it is directly supported by the mathematical or Lean files, *computed* when a supplied script reconstructs it, *author-attested* when it depends on the author’s recollection, and *inferred* when it is reconstructed from indirect agreement among sources. The main findings below rely primarily on logged, source-level, and computed support. Inferences are labelled.

The campaign boundary starts with the six-vertex census on 13 August and ends with the completion of Atlas 127 on 23 September. June–July ChatGPT work and an earlier, independent odd-cycle project are excluded. This matters because the odd-cycle theorem used in the classification was proved by a different route.

3.2 Counting rules

The reconstruction script counts typed prompts, goal-interface objectives, turns, source transitions, failed routes, commits, and retained build summaries. Provider token counters are retained separately. Codex processed tokens include cache traffic; Claude Code exposes a different decomposition; and a Lean build can occur inside a timed model turn. We therefore neither add all token counters nor sum turn durations and compilation time into one cost.

Figure 1 uses cumulative active Codex and Claude Code turn time from surviving logs as an effort proxy. Parallel turns are additive. The horizontal axis excludes human time, ChatGPT web activity, retention-deleted August Claude Code sessions, and stand-alone builds outside a logged turn. The plot is useful for chronology, not for productivity estimation. It also exposes the long interval between claimed resolution and formal status for many rows.

3.3 Selection of conversation evidence

The main text paraphrases operational exchanges in professional language. Supplementary Appendix B preserves selected source text with a stable excerpt label, date, speaker, model where known, session identifier, turn, and prompt identifier. An exchange is included only if it changed mathematical correctness, proof status, verification scope, resource policy, or research direction. This selection rule avoids treating candid wording as spectacle while allowing the reader to inspect the claims made here.

4 Findings

4.1 A formalization request triggered a reported semantic repair

Atlas 43 entered the formalization stage with a rooted sum-of-squares identity. In a 25 August Codex turn, the human author asked for Lean implementation of the unformalized proofs. Codex tested the proposed rooted semantics on the constant graphon $W \equiv 1/2$ and found that the two sides did not agree. The cause was repeated use of the same labelled edge: the informal identity treated that shared Bernoulli edge as if its conditional means multiplied independently, replacing W by W^2 .

Codex changed the semantics so repeated appearances referred to one shared Bernoulli edge variable. The Gram matrices and the 91 exact coefficient equalities remained applicable. Thus, this was a repair to the bridge from the certificate algebra to graphon density, not a replacement certificate. Codex prominently reported the error and its repair, and left Atlas 43 at **believed** until the remaining coefficient chain had been imported and checked; see Supplementary Excerpt B.1.2. It did not ask the human author to supply a revised proof.

The observable behavior supports a limited explanation. Codex treated the problem as locally repairable because the numerical and rational certificate components survived the semantic correction. Private reasoning is not available, so the paper does not claim access to why Codex chose repair over requesting a new proof. Nor did the same $W^2 = W$ mistake recur independently row by row. It was one common-framework error. Other mistakes occurred—a reversed-Jensen draft, rejected intermediate lemmas, stale status fields—but they were different failures.

This episode is more specific than the truism that checking helps. Three states diverged: the finite Gram identity could be right, its probabilistic interpretation could be wrong, and the row could remain short of an end-to-end Lean check. The Codex response separated those states rather than reporting an undifferentiated success.

4.2 Resource constraints changed the proof object

The project imposed a practical per-row envelope of roughly one hour, 16 GiB, and about 100 generated files for the compact-certificate phase. This was not merely a scheduling choice. It selected among mathematically equivalent or sufficient representations.

Atlas 43 provides the first example. A parallel Lean invocation spawned eight workers at about 13.5 GB each and exhausted machine memory. Claude Fable explicitly attributed the crash to its bare parallel build and noted that the existing verification driver was sequential (Supplementary Excerpt B.2.2). Claude Fable then compiled all 1,097 modules sequentially in 6.0 hours, with a 15.4 GiB peak (Supplementary Excerpt B.2.3). Sequential checking made that already-decomposed proof feasible, but did not make it compact.

Atlas 127 shows a stronger mathematical change. Claude Opus measured a ground-8 route intended to formalize a smoothed-Goodman theorem. An order-272 source component

used 16.20 GiB before reaching its proof checks and was killed under the cap. Claude Opus projected 800–1,200 files and 25–40 sequential hours, and explained why splitting source tables would not split the single large positive-semidefinite kernel computation (Supplementary Excerpt B.3.2).

After the author rejected a day-long compilation for one small theorem, Claude Opus weakened the target from the stronger general theorem to the Atlas 127 row itself. A direct ground-6 certificate passed the first numerical gate (Supplementary Excerpt B.3.4). A single interval still placed the Gram point on a singular face, but splitting at $p = 2/3$ produced two full-rank pieces that rationalized exactly (Supplementary Excerpt B.3.5). The resulting row used 97 files, 2617.5 seconds, and 11.90 GiB, completing the 117-row Lean status (Supplementary Excerpt B.3.6).

These episodes identify three distinct responses to size:

1. *decompose* the check and execute components sequentially;
2. *compress* a numerical search result into sparse rational certificate components; and
3. *replace* a stronger theorem by a smaller sufficient statement, then split its parameter interval to move into the interior of the exact cone.

The choice was not made by the human author alone. The author set the envelope and rejected the oversized route; Claude Opus measured the failure, identified the row-level alternative, diagnosed the singular single interval, and produced the split certificate. The resulting technical lesson is not that a small certificate is always preferable. It is that verifiability adds a second optimization problem: one must search for a proof object whose exact checking computation fits the available environment.

4.3 Formalization narrowed and reorganized the mathematics

Several informal notes stated general theorems whose full formalization was unnecessary for the finite classification. The Lean development often kept only the specialization needed for one row. This changed dependencies and, in some cases, exposed a shorter proof.

For Atlas 153, 171, 174, and 188, a target-sign observation removed self-amalgamation and Fisher-profile machinery from the route ultimately checked. Atlas 160 and 178 were reduced to smaller polynomial obligations. On the negative side, explicit tensor-product step graphons replaced an asymptotic theorem by finite rational evaluations. Atlas 127 likewise ended with the direct ground-6 certificate rather than a formal proof of the stronger smoothed-Goodman statement.

This distinction must be stated precisely. The project does not claim that Lean contains every general theorem discussed in the mathematical notes. It claims that the Lean sources prove the 117 row propositions, with reusable formal machinery for several families and row-specific endpoints elsewhere. The behavioral contrast is between informal theorem construction, which often favored conceptual scope, and formal row closure, which favored the weakest exact statement sufficient for the classification.

4.4 Codex used diverse methods in a recurring search order

The evidence does not support a claim that Codex had one fixed preferred proof method. Across the surviving sessions, it used tensor products of Turán graphons for negative rows; whiskering, self-amalgamation, cones, normalized links, and clique-tree arguments for positive families; supporting-plane, projection, and polynomial arguments for individual rows; and rooted sum-of-squares certificates for the remaining residue.

There was nevertheless a recurring order. Codex first searched for reusable families and closure principles, next developed row-specific analytic reductions, and finally built general rooted-certificate infrastructure for the unresolved cases. The failure ledger grew to 65 numbered routes and frequently distinguished a failed intermediate lemma from a counterexample to the target theorem. This hierarchy explains why the final long certificate turn was not simply an isolated burst of computation: it operated after broad families and cheaper special arguments had reduced the residue to twenty rows.

In the 25 August xhigh turn, the author asked Codex to resolve all twenty remaining cases, favoring positive proofs unless there was a strong negative indication and allowing computational exact certificates. Over about 10.5 hours, Codex built reusable certificate machinery and closed eighteen rows. It left Atlas 130 and 203 partial and described a final 42-minute solver failure as yielding no certificate or mathematical conclusion; see Supplementary Excerpt B.4.2. The combination of methodological breadth and explicitly limited failure reporting is stronger evidence than a bare count of attempted methods.

4.5 The evidence indicates an operating regime, not a winning effort tier

The productive Codex turn used xhigh reasoning effort, while eleven live ultra turns on the second device totalled 4.16 hours and closed no row. The only high Codex proving turn in the relevant comparison was a 28-minute failed audit. A simple slogan that higher reasoning effort is worse would be unsupported. The sessions differ in task, duration, context, continuity, model state, and available tooling.

What can be said is narrower and more useful. Closure occurred in a regime combining four features: a long uninterrupted turn, a sharply stated finite-row objective, permission to build exact computational machinery, and a residue already prepared by previous analytic work. The xhigh label is one part of that configuration. The ultra attempts were shorter and fragmented; the high attempt was both short and aimed at a different audit problem. Thus, the history motivates a hypothesis that there is a suitable allocation of continuity, tool-building effort, and task scope. It does not estimate the independent effect of a reasoning-effort setting.

5 What the quantitative traces do and do not show

The largest counters in Table 2 are not comparable measures of mathematical work. The workstation Codex total includes 1,868,302,720 cached input tokens. The second-device figure uses the provider’s processed-token counter over primary threads. Claude Code separates input, cache traffic, and output differently. Replayed forks can also repeat large contexts. These quantities are useful for auditing the scale of interaction and the danger of treating context replay as fresh reasoning; they are not prices or additive compute estimates.

The same caution applies to time. Twenty retained bounded row builds total 11.69 hours, with a largest peak of 11.91 GiB, while the separate Atlas 43 check took 6.0 hours. Some compilation happened inside Claude Code or Codex turns. Adding model-turn and build time would double count those intervals. The surviving sources also omit human reading time, several deleted sessions, ChatGPT web duration, energy use, and the unavailable full-build directory.

The author reports a ChatGPT Pro subscription price of KRW 159,000 per month and Claude Max charges of USD 96.76 on 21 July and USD 110 on 21 September. These subscriptions covered activity outside the defined campaign. They are reported for transparency but are not assigned as a project cost.

Table 2: Observable campaign quantities. Token counters retain each provider’s semantics and are not added across providers.

quantity	value	definition
Human prompts in the workstation appendix	86	typed prompts to Codex or Claude Code
Codex goal objectives	10	goal-interface objectives
Workstation Codex worker turns	107	all extracted Codex worker turns
Second-device primary live turns	40	non-replayed primary Codex turns
Workstation Codex processed tokens	1,912,053,938	provider counter; includes cache traffic
Second-device Codex processed tokens	669,491,502	provider counter; primary threads
Claude Code input tokens	283,520,656	input plus cache read and cache creation
Claude Code output tokens	1,333,543	surviving workstation sessions
ChatGPT Pro conversations	16	campaign plus satellite conversations
Automatic approval reviews	45	29 workstation plus 16 second-device requests
Numbered failed routes	65	entries in the failure ledger
Bounded sequential row builds	20	11.69 h combined; peak 11.91 GiB
Atlas 43 sequential chain	1,097	6.0 h; peak 15.4 GiB

6 Reusable artifacts and epistemic drift

Two shared files were especially important. The catalogue maintained one row per graph with mathematical status, proof reason, witness, and Lean state. The failed-route ledger preserved negative knowledge, including exact counterexamples to intermediate lemmas for rows whose target theorem was eventually positive. Together they allowed later Codex and Claude Code sessions to begin from more than a prose handoff.

Their benefit is plausible but not experimentally isolated. The human author asked that these files be maintained; the systems did not independently choose a modular division of labor. Moreover, the catalogue itself drifted. Narrative sections and architecture notes lagged behind source changes, a count module required correction, and Atlas 43 and 196 were marked verified before the later measured sequential check. A shared catalogue was therefore useful working memory and also a boundary at which confidence labels could get ahead of executable evidence.

The final theorem artifacts are easier to reproduce than their discovery. Exact rational witnesses, certificate checkers, Lean files, twenty row-build summaries, and the Atlas 43 compilation summary survive. Numerical semidefinite workspaces, some deleted Claude Code transcripts, private model reasoning, and the full-build directory do not. A future reader can inspect many final proof obligations without being able to replay every search choice that produced them.

7 A prespecified clean-rerun protocol

The historical material suggests that structured handoffs, exact-certificate infrastructure, and a resource-aware objective may help close a finite table. It cannot test that claim. We propose the following experiment for a future paper.

7.1 Snapshot and tasks

Freeze a repository snapshot immediately before the twenty-row residual campaign. Remove later proofs, generated certificates, completion-specific notes, and commit messages while retaining the common mathematical foundation available at that date. Package each residual row with the same statement, permitted dependencies, and checker interface. Because rows differ greatly in difficulty, use the entire twenty-row bundle as the primary task rather than treating rows as exchangeable independent samples.

7.2 Matched conditions

Run at least five independent trials per condition with the same Codex model version, effort setting, token allowance, wall time, machine class, and base snapshot. Compare:

1. a direct instruction to solve the residual rows; and
2. a structured workflow supplying the catalogue, failed-route format, exact-certificate tools, explicit resource envelope, and a requirement to separate conjectured, exact-checked, Lean-checked, and fresh-built states.

The human may respond only through a prespecified heartbeat message and may not contribute new mathematics during a trial. Any crash-recovery policy and permission decision must be fixed in advance.

7.3 Endpoints

The primary endpoint is the number of rows that pass a fresh Lean build under the stated envelope. Secondary endpoints are time to first verified row, number of exact certificates produced, false positive status transitions, maximum memory, generated source size, failed routes revisited, and the gap between claimed proof and fresh compilation. An independent script should assign statuses from artifacts rather than from the model’s final prose.

A separate factorial experiment is needed to study reasoning effort. It should hold the model, snapshot, prompt, tools, wall time, and trial count fixed while varying only high, xhigh, and ultra. Without that design, the current campaign should not be cited as evidence that one effort setting is better.

7.4 Threats to the rerun

Public availability of the completed repository may contaminate later model versions. A private isomorphic benchmark or newly generated family would reduce direct memorization. Trials may still share provider-side state that is unavailable to the investigator. Lean success also favors work already close to the supplied library. The protocol therefore tests a workflow on a bounded formal task; it does not measure mathematical research ability in general.

8 Limitations

This study has one campaign, one mathematical domain, and no randomized assignment. The human author changed prompts after observing failures, so interventions are endogenous to progress. Models and interfaces changed across sessions. Several August Claude Code transcripts were removed by retention, the ChatGPT export omits some context and files, and private model reasoning is unavailable. The unavailable full-build directory prevents an independent reconstruction of that run’s environment and timing from this checkout.

Status labels are historical claims, not substitutes for compilation. In particular, the completion curve follows source state and includes the period when Atlas 43 and 196 were marked verified before the retained sequential audit. Token counters have provider-specific meanings, and wall time is incomplete. Subscription payments cannot be allocated to this project.

Finally, quotations were selected because they changed correctness, status, scope, resources, or direction. They are not a random sample of interaction quality. The appendix supplies enough surrounding context to audit the claims made here, while the full indexes remain necessary for alternative analyses.

9 Conclusion

The campaign shows a progression from broad mathematical search to compact, machine-checkable row proofs, but its most useful lessons lie in contrasts. Codex repaired a semantic error while preserving the algebraic certificate; Claude Code changed an infeasible strong theorem into a feasible row-level certificate; Lean formalization repeatedly selected smaller sufficient statements; and Codex moved through several analytic families before building general exact-certificate machinery. The successful xhigh run combined continuity, a finite objective, and reusable tools, so its outcome cannot be attributed to effort level alone.

For future AI-assisted mathematics, the checkable endpoint should be defined before the search begins, and proof-object resources should be treated as part of proof design. Claims of workflow benefit or model superiority require the clean matched reruns that this historical campaign did not contain.

AI assistance and conflict disclosure. OpenAI Codex (GPT-5) drafted this article from the project sources and wrote the scripts that regenerate its aggregate counts and curated supplementary excerpts. Codex is an OpenAI system, and OpenAI systems are among the subjects of this case study. The authors remain responsible for source selection, interpretation, and the submitted text.

A Audit definitions and source gaps

A.1 Status definitions

Claimed resolved

A surviving source asserts either a positive proof or an exact negative witness for the row.

Exact-checked

A rational checker verifies the finite certificate or witness arithmetic independently of a numerical solver.

Lean-marked verified

The historical Lean source status is `.verified`; this is the state used by Figure 1.

Fresh-built

The relevant source was compiled from a specified source state in a retained sequential run.

Full-build attestation

The author reports checking the entire development on another device, but the run directory is unavailable here.

A.2 Known gaps

The supplied checkout lacks the referenced full-run directory, `lean/verification_runs/full-20260923-164651/`; a handoff named `HANDOFF_ATLAS_188_171_174_153.md`; and `lean/docs/TODO.md`. A search of the reachable git history did not recover the latter two paths. The workstation raw-log backup is absent, although its extracted index and appendices survive. The second-device backup survives. These gaps chiefly constrain exact discovery provenance and full-build timing, not the inspectability of the final row sources.

B Curated chronological excerpts

The selection rule includes exchanges that changed mathematical correctness, proof status, verification scope, resource policy, or research direction. Wording and typographical errors are retained. Export-only HTML spacing markers are omitted, and mathematical Unicode glyphs receive equivalent LaTeX typesetting. Ellipses are not inserted inside an excerpt.

B.1 The Atlas 43 semantic repair

Excerpt B.1.1

25 August 2026, 13:21 KST; Codex gpt-5.6-sol, xhigh; session 01a03722, turn 1; prompt A.38; speaker: human author.

Thank you. There are lots of unverified (unformalised) proofs now.
Would you please work on implementing the lean proof? You may find the lean proof existing first to see how the proofs are organised.

Excerpt B.1.2

25 August 2026, 14:04 KST; same session and turn; speaker: Codex.

Key finding: the incoming SOS identity was unsound as written. Shared labelled edges were treated as though $W(x,y)^2 = W(x,y)$. For $W \equiv 1/2$, the claimed sides evaluate to 0.015625 and approximately 0.0038337 . I repaired the formulation using shared Bernoulli edge variables.

Atlas 43 remains marked ‘believed’: the remaining task is importing the large rational certificate and verifying its 91 coefficient equations in Lean, then connecting that result to `Graph043.status`.

B.2 Memory failure and sequential checking

Excerpt B.2.1

22 September 2026, 13:27 KST; Claude Code claude-fable-5-1, high; session d19b441c, turn 6; prompt A.74; speaker: human author.

We ran out of memory and whole computer stopped. Was the previous verification verified module in parallel? If so, we might try sequential build, with at most one file to build. Also, consider 16 or 32 GB RAM usage upper limit.

Excerpt B.2.2

22 September 2026, 14:30 KST; same session and turn; speaker: Claude Fable.

1. ****The crash was mine.**** I ran a bare ‘lake build’, which spawned 8 workers at ~13.5 GB each. Codex’s runs were always sequential via ‘experiments/verify_lean_modules.py’. I’ve saved that as a memory rule and noted it in the plan and the record so it can’t recur.

Excerpt B.2.3

22 September 2026, 22:07 KST; same session, turn 8; speaker: Claude Fable.

****Atlas 43 and 196 are ✓ Verified again****, on real evidence this time:

- All 1,097 modules of the Atlas 43 chain compiled sequentially and correctly (6.0 h total, mean 20 s/module, peak 15.4 GiB on ‘CoresRows56’, which needs ‘-M 16384’; nothing needed more than 16 GiB).
- ‘Graph043.status’ and ‘Graph196.status’ are the real proofs and audit to ‘[proptext, Classical.choice, Quot.sound]’.
- Table: ****112 verified, 5 believed**** (127, 153, 171, 174, 188), 0 unresolved. The row cells cite the record and the six-hour cost.

B.3 Replacing the Atlas 127 proof object

Excerpt B.3.1

23 September 2026, 08:02 KST; Claude Code claude-opus-5, *xhigh*; session 272ea172, turn 2; prompt A.83; speaker: human author.

Okay, now we have huge one problem, 127. I don’t think current proof direction is feasible. Would you please guess how many proof files it need, and how long it would take to compile them without actually trying entire compilation? Please make sure to delete generated files after this check.

Excerpt B.3.2

23 September 2026, 11:44 KST; same session and turn; speaker: Claude Opus.

Atlas 127: not feasible on the current route

****The blocker is not total cost -- it’s that one block cannot be done at all.****

I synthesized a block module at Atlas 127’s largest order using the real generator’s own emitters (‘matrix’, ‘packed_product’), with true statements so the kernel checks actually run. Two measurements, each a single sequential module under the standard 16 GiB watchdog:

synthetic block	source	result
order 102 (matches today’s largest)	2.1 MB	Data 133.9 s / 6.54 GiB, checks 114.9 s / 9.13 GiB -- **passes**
order 272 (Atlas 127’s largest)	15.0 MB	Data alone: 2,259 s, **16.20 GiB -- killed by the memory cap, never reached the checks**

The order-102 run reproduces the real ‘Block01’ (303 s / 11.90 GiB) closely, so the harness is faithful and if anything optimistic.

Atlas 127 has **11** blocks of order 252-272 (the $m=4$ flag types), carrying 776,564 of its 1,431,052 Gram entries. Each one exceeds the per-row envelope on memory before any proof obligation is attempted. The ‘Data’ tables can be split across modules, but a congruence PSD proof over a 272×272 matrix is a single kernel evaluation and cannot be split -- so no amount of file-splitting rescues these 11 blocks. The other 173 blocks (order ≤ 109) are individually fine.

Files: roughly 800-1,200.

- ~420 for the 184 PSD blocks (133 order-64 blocks at 2 files each; the 11 large blocks needing 5-6 each even to hold their data)
- ~150 for the slack tables (10-20 M integers $\Rightarrow$ ~375 MB of source at ~2.5 MB/module)
- ~250 for the 12,346 slack inequality checks, chunked as the S4 contraction checks are
- plus coloring and assembly

Time: roughly 25-40 hours sequential, from the measured 120-200 s per MB of generated source across ~660 MB total. For scale, Atlas 43’s retired 1,097-module chain took 6 h.

That is 25-40 $\times$ the 3,600 s per-row envelope, in ~10 $\times$ the file budget.

Excerpt B.3.3

23 September 2026, 11:44 KST; same session, turn 3; prompt A.84; speaker: human author.

I see. One small theorem taking more than one day to compile is a bit overkill... Can you think of some directions that we can reduce the size of proof and get in feasible size?

Excerpt B.3.4

23 September 2026, 12:00 KST; same session and turn; speaker: Claude Opus.

The decisive direction works -- at least at its first gate. Atlas 127’s own catalogue bound is feasible in the ground-6 four-root cone that just verified sixteen rows:

“

```
atlas=127 interval=[1/2,1] equations=407 rank=407 augmented=407
status=optimal max_residual=1.99e-06
G0_4 rank 57/64 G0_31 rank 101/104 G0_22 rank 65/68 G0_211 60/60 ...
```

“

Same signature as every row that later rationalized successfully. No files were written (no ‘-save’).

Why this is the answer

The 25-40 hour estimate is for formalizing the **smoothed Goodman theorem**, ‘ $t(\theta_{1,2,4},W) \geq (2p-1)t(C_5,W)$ ’ -- a standalone result that is

strictly stronger than what Atlas 127 needs. Atlas 127 only needs ' $\Phi_{127}(p) \leq t(127, W)$ ' for ' $p \geq 1/2$ '. The plan's row 20 says exactly this ("First seek a smaller direct certificate for the catalogue bound"), and as far as I can tell it was never tried -- 127 appears nowhere in the interval-SOS table or the campaign records.

If the rationalization goes through, Atlas 127 becomes a single-interval row like 118/122/124: ****~89 files, ~25 minutes, ~9 GiB**** -- inside the normal envelope, with no new machinery at all. The whole downstream pipeline accepts ' $[1/2, 1]$ ' unchanged.

Two honest caveats. Numerical feasibility is not exact feasibility; the double root of Φ_{127} at $p = 1/2$ makes the bound tangent to second order there, which is harder to round than the simple roots the other rows had. And this verifies **the row**, not **the theorem** -- the smoothed Goodman paper would remain unformalized.

Excerpt B.3.5

23 September 2026, 12:49 KST; same session, turn 4; response to prompt A.85; speaker: Claude Opus.

Rationalization works. Atlas 127 is an ordinary two-piece row.

Both gates pass, with margins better than several rows already verified.

| | ' $[1/2, 2/3]$ ' | ' $[2/3, 1]$ ' |

|--|--|--|

| margin SDP | CLARABEL 'optimal', $2.24e-05$, all blocks full rank |
CLARABEL 'optimal', $1.47e-04$, all full rank |

| exact interval SOS | ****all 407 equations verified****, min DD margin ****0.99585**** | ****all 407 verified****, min DD margin ****0.99972**** |

| compact face | min DD margin ****0.9999985****, denominator ****18 digits**** |
| min DD margin ****0.9999994****, denominator ****17 digits**** |

| independent audit | passed -- 407 equations, 10 positive matrices |
passed |

| block orders | 57, 101, 65, 60, 12 / 32, 52, 34, 30, 6 | 50, 98, 62,
60, 12 / 23, 48, 30, 30, 6 |

Both ' ϕ ' fields back-substitute exactly to $\Phi_{127} = 8p^5 - 16p^4 + 14p^3 - 6p^2 + p$, and the pieces meet at $2/3$ and cover the full admissible ' $[1/2, 1]$ ' for $\chi = 3$. The block orders are identical to Atlas 151's, and the denominators are ****shorter**** than all four rows I verified today (188: 19 digits, 171: 22, 174: 23, 153: 19).

So Atlas 127 projects as a standard 'Lower'/'Upper' row -- ****~97 files, ~40 minutes, ~11.9 GiB**** -- needing no new machinery, no low module, and no changes to the pipeline. That replaces the ~1,000-file, 25-40-hour ground-8 route.

****The single interval was the whole problem.**** ' $[1/2, 1]$ ' has maximum margin exactly 0 (CLARABEL $-6.9e-08$, SCS $+3.5e-08$, ranks collapsing to 32/50 and 59/98) -- every feasible Gram is singular there, so the Cholesky in the rationalizer cannot round it. Split at $2/3$, both

pieces are strictly interior at full rank. My earlier guess that Φ_{127} 's double root at $p = 1/2$ was to blame was wrong: Atlas 151 has the same $(2p-1)^2$ tangency and rationalized fine, because it was split. This also retroactively explains the handoff's remark that the August campaign split every row at $2/3$ "possibly only for conditioning" -- for 127 the split is not cosmetic, it is the difference between a degenerate and a non-degenerate cone.

Excerpt B.3.6

23 September 2026, 13:44 KST; same session, turn 5; response to prompt A.86; speaker: Claude Opus.

Atlas 127 is verified and committed ('87a8e439f').

| | |

|--|--|

| fresh sequential build | **2,609.6 s**, 97 files, peak **11.90 GiB**
|

| installed recheck | 7.8 s -- combined **2,617.5 s** |

| axioms | 'proptext', 'Classical.choice', 'Quot.sound' only, in both
the fresh catalogue compile and the recheck |

| envelope | $97 \leq 100$ files, $2,617 \leq 3,600$ s, $11.90 \leq 16$ GiB |

****All 117 classification rows are now Lean-verified -- 0 believed, 0 unresolved.****

It went through as an ordinary two-piece 'Lower'/'Upper' row, no new machinery, no low module: the certificates rationalized as predicted, the compact denominators came out at 18 and 17 digits (the shortest of the recent rows), and the build landed within 2 % of Atlas 151, whose block orders are identical.

B.4 The long Codex certificate turn

Excerpt B.4.1

25 August 2026, 01:34 KST; Codex gpt-5.6-sol, xhigh; session 01a0349f, turn 0; prompt D.87; speaker: human author.

/goal Would you please read the
GRAPH_CLASSIFICATION_UP_TO_6_VERTICES.md, and see some related
files if you need, and try to prove remaining instances that they
satisfy the conjectured lower bound? I think most of the instances are
heavily tested for counterexample, so unless you have strong clue that
it is negative, try to prove in the positive direction. Work until you
prove all remaining 20 instances (14 open + 6 partial) as positive or
negative.

Excerpt B.4.2

25 August 2026, 12:03 KST; same session and turn; speaker: Codex.

Outcome:

- 18 of the 20 remaining cases are now proved positive.
- Atlas 130 and 203 remain partial:
- Atlas 130: proved on $\backslash(p\text{in}[2/3,1]\backslash)$; $\backslash([1/2,2/3]\backslash)$ unresolved.
- Atlas 203: proved on $\backslash(p\text{in}[3/4,1]\backslash)$; $\backslash([2/3,3/4]\backslash)$ unresolved.
- No negative instance was found.
- The final Atlas 203 CVXOPT job ran about 42 minutes, failed at the solver level, and produced no certificate or mathematical conclusion.

References

- [1] Authors. A complete six-vertex classification for a chromatic graphon lower bound. Companion manuscript, 2026.
- [2] Emily First, Markus N. Rabe, Talia Ringer, and Yuriy Brun. Baldur: Whole-proof generation and repair with large language models. In *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, pages 1229–1241, 2023.
- [3] Stanislas Polu and Ilya Sutskever. Generative language modeling for automated theorem proving. *arXiv preprint arXiv:2009.03393*, 2020.
- [4] Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, Matej Balog, M. Pawan Kumar, Emilien Dupont, Francisco J. R. Ruiz, Jordan S. Ellenberg, Pengming Wang, Omar Fawzi, Pushmeet Kohli, and Alhussein Fawzi. Mathematical discoveries from program search with large language models. *Nature*, 625:468–475, 2024.
- [5] Trieu H. Trinh, Yuhuai Wu, Quoc V. Le, He He, and Thang Luong. Solving olympiad geometry without human demonstrations. *Nature*, 625:476–482, 2024.
- [6] Kaiyu Yang, Aidan M. Swope, Alex Gu, Rahul Chalamala, Peiyang Song, Shixing Yu, Saad Godil, Ryan Prenger, and Anima Anandkumar. LeanDojo: Theorem proving with retrieval-augmented language models. In *Advances in Neural Information Processing Systems*, volume 36, 2023.